

PostgreSQL

ITWS 2110 & CSCI 2960
Web Systems Development
Fall 2015

Ahmed Eleish

Agenda

- Intro
- Postgres
- Object-relational DBMS
- Features of PostgreSQL
 - Inheritance
 - Datatypes
 - GiST
 - Compliance
 - MVCC
 - Full text search
 - Extensibility
 - Community
- PHP and PostgreSQL

Database

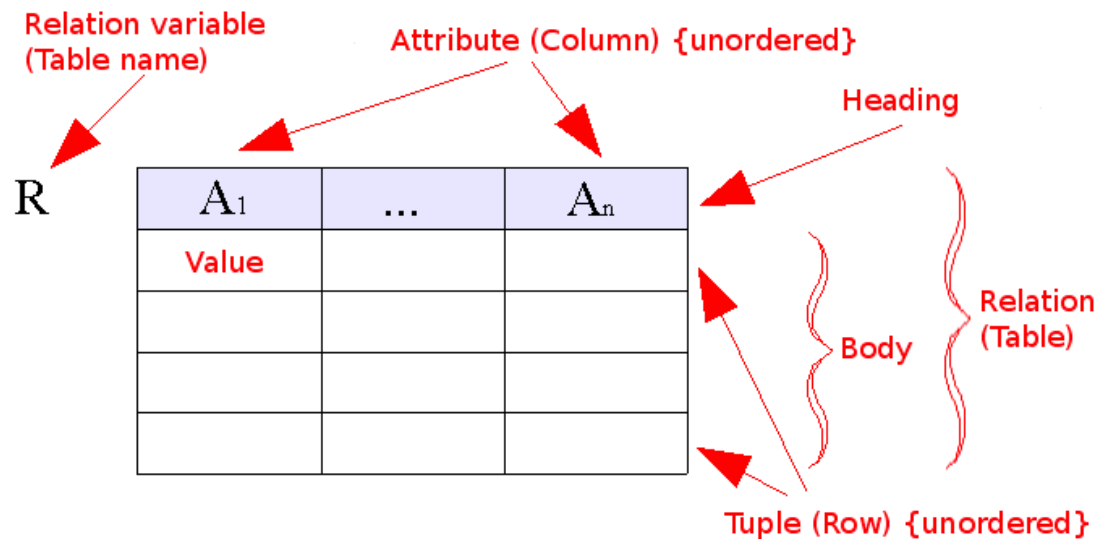

- An organized collection of data
- Has a database model
- Is managed by a Database Management System

Database Management System

- Application that manages the creation, querying, update, and administration of databases
- Supports particular logical and physical data models
- Allows users to create/manipulate data/information while taking care of the underlying data structures

Relational Database Management Systems

- Popular since the 1980s
- MySQL, PostgreSQL, Oracle, MS SQL Server, IBM DB2
- Implement the **Relational Model** proposed by E.F. Codd in 1969



Relational Database Management Systems

- Tables, keys, and constraints
- Compact and clean
- Mathematical base allows deep analysis of data

Employees Table

| id | First Name | Last Name | Department |
|-----|------------|-----------|------------|
| 001 | Jane | Doe | 00 |
| 002 | John | Smith | 01 |
| 003 | John | Doe | 01 |

Key=01

Departments Table

| id | Name |
|----|----------------|
| 00 | Management |
| 01 | Not Management |

Structured Query Language (SQL)

- Programming language for RDBMS
- Originally based on relational algebra & relational calculus
- Data Definition Language (DDL):
 - CREATE, ALTER, DROP
- Data Manipulation Language (DML):
 - SELECT..FROM, INSERT..INTO, UPDATE..SET, DELETE..FROM
- Date Control Language (DCL):
 - GRANT, REVOKE

ACID-Compliance

- Atomicity
 - “All or nothing”: if part of the transaction fails the entire transaction fails
- Consistency
 - Write operations will bring the DB to a valid state, according to constraints, cascades, and triggers
- Isolation
 - Serializability, result of concurrent execution of transactions as if executed serially
- Durability
 - Committed transactions will not be lost due to hardware/software failure

PostgreSQL

- Object-relational database system
- Runs on all major OS
- ACID-compliant
- Programming interfaces for Java, .NET, Perl, Python, Ruby, C/C++, and its own PL/pgSQL
- Enterprise class, highly scalable

PostgreSQL

- Evolved from project Ingres at UC, Berkeley
- Originally titled Postgres
- Renamed PostgreSQL in '96
- Developed by PostgreSQL Global Development Group
- **Free and open-source**

Object-relational DBMS

- Database model is hybrid of relational model and **object-oriented model**

Object 1: Employees


| | |
|------------|--|
| id | |
| First Name | |
| Last Name | |
| Department | |

Object 1 instance

| |
|------|
| 001 |
| Jane |
| Doe |
| 00 |

Object 2: Departments

| | |
|------------|--|
| id | |
| First Name | |
| Last Name | |
| Department | |



Why Object-relational

- Relational Database
 - Mathematical relation allows set theory techniques in data analysis
 - Impedance mismatch between application and DB
- Object-oriented Database
 - Allows sets, lists, user-defined datatypes and nested objects
 - No mathematical base for deep analysis
- ORDBMS bridge the gap
 - Complex data, nested objects, type inheritance
 - Mathematical base for analysis

PostgreSQL: Type Inheritance

- Tables can inherit from parent tables
- Data in child tables will appear as if they exist in parent tables
- Can be used for data partitioning
 - e.g. Separate tables for monthly sales data
 - When querying date range only relevant tables are scanned
 - Smaller tables with smaller indexes mean faster searches
 - Easier maintenance
- Has shortcomings regarding unique constraints and foreign keys

PostgreSQL: Supported Datatypes

- Native datatypes:
 - Boolean, Character, Binary, Date/Time, bigint, integer, double precision, XML, JSON, IPv4, IPv6 addresses, MAC addresses
- User-defined datatypes:
 - Complex objects, domains, ranges, geometric types such as points, lineStrings, Polygons,...
 - Access methods for these types need to be defined

Generalized Search Tree (GiST)

- Disk-based indexing system, generalized form of the B+ tree
- Allows creation of custom datatypes with indexed access methods
- Can be used to implement B+ Trees, R-Trees, among others
- Includes support for variable-length keys and composite keys
- GiST is extensible

Generalized Search Tree (GiST)

- Can support nearest neighbor search (KNN-GIST)
- Can be used to find similar words or close objects or locations within geospatial data
 - This makes it ideal for GIS databases

PostGIS

- Adds support for geographic objects to PostgreSQL
- Follows the Simple Features spec. from the Open Geographic Consortium
- Features
 - Geometry types for Points, LineStrings, Polygons, MultiPoints,...
 - Spatial predicates
 - Geospatial operators for measurements and set operations

PostgreSQL: Compliance

- Implements majority of SQL:2011
 - Makes migration of data to other systems easy
- Fully ACID-compliant
 - High reliability and emphasis on data integrity
 - Implements MVCC to manage concurrency

Multiversion Concurrency Control

- Every transaction gets a “snapshot” of the database
- Uses timestamps and incrementing transaction IDs
- A transaction sees only transactions completed before it started

Multiversion Concurrency Control

- On query start PostgreSQL records the transaction counter
- Visible rows must have a creation transaction id that:
 - Is a committed transaction
 - is less than the transaction counter recorded at query start
and
 - was not in-process at query start
- Visible rows must also have an expire transaction id that:
 - is blank or aborted or
 - is greater than the id stored at query start
or
 - was in-process at query start

Multiversion Concurrency Control

Cre 30
Exp

Create

Cre 40
Exp 47

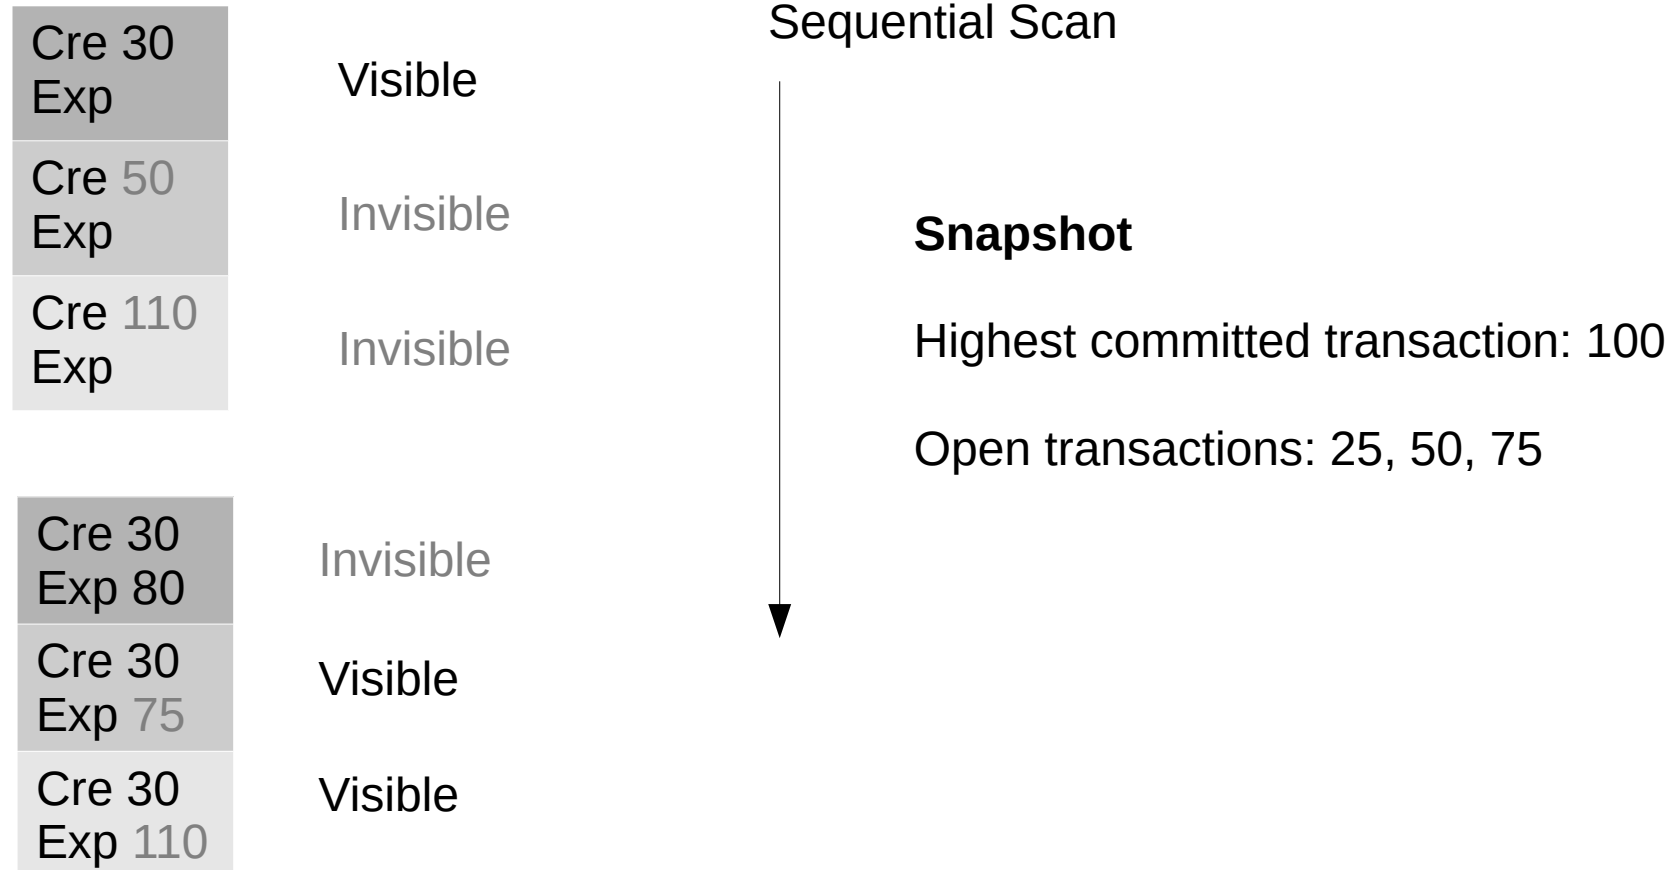
Delete

Cre 64
Exp 78

Cre 78
Exp

Update

Multiversion Concurrency Control



PostgreSQL: Full Text Search

- Problems with regular text search in DB systems:
 - No linguistic support
 - Search results are not ranked
 - Slow without index support
- In PostgreSQL:
 - Documents are preprocessed by:
 - Parsing documents into tokens
 - Converting tokens into lexemes using *dictionaries*
 - Storing preprocessed documents

PostgreSQL: Full Text Search

- Dictionaries:
 - Programs that convert tokens into lexemes
 - Normalize words so that derived forms of the same word will match
 - PostgreSQL provides predefined dictionaries for many languages
 - Templates are available to create new dictionaries

PostgreSQL: Extensibility

- Open-source with permissive free-software license
- Programming interfaces for C/C++, Java, .NET, Perl, Python, ODBC,...
- Stored procedures can be written in over a dozen languages and loaded as libraries
- Many third-party tools for designing and managing the system

PostgreSQL: Extensibility

- GiST enables the creation of new indexing methods allowing users to specify..
 - What to store
 - How to store it
 - and
 - Define new ways of search it

PostgreSQL: Community

- PostgreSQL Global Development Group
 - Diverse group of companies and individual contributors
- Active community including mailing lists and IRC
- Commercial support options also available

When to use PostgreSQL

- Where standards and compliance to SQL are expected
- If data integrity and reliability are the top priority
- When complex, custom procedures are needed
- Where complex database designs are required
- If an enterprise class DBMS is needed with low costs

When not to use PostgreSQL

- In systems with a simple structure
- Where *read* performance is the top requirement
- Where database and system administration experience may be lacking

PHP and PostgreSQL

- PHP includes an API for connecting to, and manipulating data in a PostgreSQL database
- PostgreSQL 6.5 or later
- Connecting to DB
 - `pg_connect`, `pg_close`
- Manipulating data
 - `pg_query`, `pg_update`
- Prepared Statements
 - `pg_prepare`, `pg_execute`

PHP and PostgreSQL

- PHP also includes PDO_PGSQL which implements PDO to enable access to PostgreSQL
- Functions
 - PDO_PGSQL DSN – Connecting to database
 - PDO::pgsqlCopyToArray – Copying data from db to PHP